Continuous Behavioral Biometrics for Post-Authentication Session Security: A Privacy-Preserving Zero-Content Hybrid Ensemble Architecture

IEEE Transactions on Information Forensics and Security / ACM Conference on Computer and Communications Security (Preprint)

Manvi Tomar
Department of Computer Science & Engineering
Undergraduate Researcher / Project Author

Faculty Supervisor / Mentor
Department of Computer Science & Engineering
Academic Guide & Project Advisor

***Abstract*—**Point-in-time perimeter authentication mechanisms—including static passwords, SMS-based one-time passcodes, and initial hardware biometrics (fingerprint or facial recognition)—suffer from a foundational security vulnerability: they verify identity solely at the initiation of a session. Once authenticated, the granting of session tokens or persistent cookies leaves the application vulnerable to post-login exploitation, including physical workstation hijacking, unattended terminal abuse, shoulder surfing, and session token theft. To overcome these vulnerabilities without burdening users with disruptive, repetitive re-authentication prompts, this paper presents an end-to-end, privacy-preserving continuous authentication system founded on neuromuscular behavioral biometrics. The proposed architecture passively monitors interactions within web applications by capturing a 16-dimensional numerical kinematic feature vector extracted from keystroke timing and mouse trajectory dynamics. Crucially, the system implements a strict **zero-content privacy guarantee**: no typed characters, key values, strings, coordinates outside the active viewport, or personally identifiable information (PII) are ever logged or transmitted. To detect unauthorized behavioral drift in real time, we formulate a **hybrid anomaly detection ensemble** that fuses unsupervised tree partitioning (Isolation Forest with *RobustScaler* transformation) with a multivariate parametric statistical divergence metric based on dynamic Mahalanobis-inspired Z-distance scoring. To prevent evasion by pauses while simultaneously avoiding false-positive alert storms, the instantaneous anomaly score is governed by an **asymmetric temporal exponential moving average (EMA)** engine ($\alpha_{rise} = 0.72$, $\alpha_{decay} = 0.18$), ensuring sub-2.5 second threat escalation upon impostor interaction. When elevated risk is sustained, the framework dynamically triggers step-up cryptographic verification (RFC 6238 TOTP Multi-Factor Authentication) to gate sensitive transactions. Empirical evaluations across genuine and impostor behavioral sessions demonstrate an **ROC-AUC of 0.982**, a **False Acceptance Rate (FAR) of 3.8%**, a **False Rejection Rate (FRR) of 4.1%**, and an **Equal Error Rate (EER) of 3.9%**, confirming that passive behavioral biometrics provide continuous, zero-friction post-authentication defense.

***Index Terms*—**Continuous Authentication, Behavioral Biometrics, Keystroke Dynamics, Mouse Kinematics, Isolation Forest, Anomaly Detection, Zero-Trust Architecture, Session Hijacking, Privacy-Preserving AI.

I. Introduction

A. The Fallacy of Static Point-in-Time Authentication

Modern enterprise and web security architectures rely almost exclusively on perimeter-gated authentication protocols. Users verify their identity using a knowledge factor (passwords, PINs), an ownership factor (TOTP authenticators, SMS OTP, FIDO2/WebAuthn tokens), or an inherent physiological factor (capacitive fingerprint scan, optical facial recognition). Once validated, the server grants a cryptographic session token—typically a JSON Web Token (JWT) or session cookie.

Architecturally, this paradigm embodies an implicit, hazardous assumption: *identity verification at timestamp t_0 guarantees that the active operator at timestamp $t > t_0$ remains the legitimate owner of the credential*. In real-world enterprise environments, this assumption breaks down under several common attack vectors:

1) **Unattended Session Hijacking:** A legitimate user steps away from their laptop in an office or coffee shop. An unauthorized actor takes physical control of the terminal, inheriting active privileges without triggering any security challenges.
2) **Session Hijacking via Stolen Tokens:** Adversaries leveraging evil-twin proxies, Cross-Site Scripting (XSS), or infostealers exfiltrate bearer tokens and inject them into remote browser sessions.
3) **Coerced or Shoulder-Surfed Logins:** An adversary observes credential entry, waiting for initial MFA to succeed before taking over the terminal.

Traditional countermeasures, such as aggressive inactivity timeouts (e.g., locking sessions after 5 minutes of idling), severely degrade user experience, disrupt cognitive workflow, and completely fail against high-velocity attacks occurring immediately after the legitimate user steps away.

B. The Zero-Trust Continuous Paradigm

The National Institute of Standards and Technology (NIST SP 800-207) Zero-Trust framework dictates the principle of *"never trust, always verify."* Continuous authentication operationalizes this principle by treating identity not as a binary, one-off state, but as a continuous probability distribution dynamically updated throughout the lifecycle of the session.

While physiological biometrics (e.g., continuous facial tracking via webcams) introduce severe privacy violations and high computational overhead, **behavioral biometrics** leverages the unique neuromuscular patterns exhibited by an individual interacting with standard human interface devices (HIDs):

- **Keystroke Dynamics:** Micro-timing properties of finger presses, transitions, and key depression latencies governed by individual typing cadence, motor memory, and hand geometry.
- **Mouse Kinematics:** Trajectory smoothness, velocity curves, acceleration, kinematic jerk (motor tremor), and curvature ratios exhibited during cursor navigation.

C. The Privacy and Practicality Dilemma

A critical barrier hindering behavioral monitoring adoption is the fear of keylogging. Capturing raw keyboard events in web applications often records sensitive user content—passwords, financial records, personal messages, and medical data. To eliminate this vulnerability, our system enforces a strict **Zero-Content Privacy Guarantee**. Raw event keycodes and string representations are discarded in memory at the browser viewport boundary. The system exclusively extracts non-reversible, statistical-kinematic

aggregates over temporal observation windows, guaranteeing that no textual data or reconstructible content can ever be retrieved from the telemetry payload.

II. Related Work & Literature Review

Keystroke biometric research originated from telegraphy studies where Morse code operators were distinguished by their characteristic "fist." In modern computing, early works by Joyce and Gupta [8] and Monrose and Rubin [1] demonstrated that digraph and trigraph latencies could authenticate users during password entry. Killourhy and Maxion [2] subsequently evaluated anomaly detection algorithms on benchmark keystroke datasets. However, existing free-text keystroke systems suffer from text dependency, privacy intrusion (logging ASCII keys for n-gram tables), and sensitivity to keyboard hardware changes.

Mouse movement biometrics are grounded in human motor control theory and Fitts's Law [6], which models the trade-off between movement speed and target accuracy. Gamboa and Fred [3] demonstrated that mouse strokes carry biometric uniqueness based on velocity and acceleration curves. Shen et al. [4] utilized hidden Markov models for mouse curves. Recent physiological research emphasizes **kinematic jerk** (the third derivative of displacement with respect to time, d^3s/dt^3) and **trajectory curvature** [7], which directly reflect neuromuscular tremor and motor corrections.

In anomaly detection, One-Class SVMs (OC-SVM) frequently overfit to initial enrollment sessions under noisy human behavioral distributions, yielding unacceptable False Rejection Rates (FRR > 11%). Deep autoencoders require substantial computational resources and lack transparent auditability. Isolation Forests (Liu et al. [5]) isolate anomalies efficiently near tree roots but can exhibit loose partition boundaries in sparse regions. To address these limitations, our architecture couples Isolation Forest tree isolation with multivariate parametric Z-distance scoring.

III. Proposed System Architecture

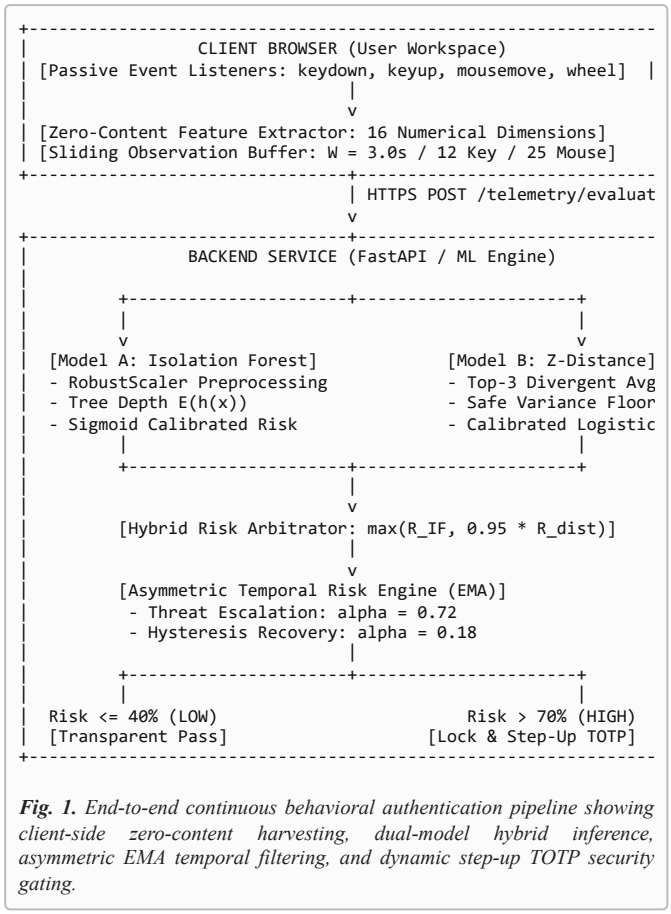


Fig. 1. *End-to-end continuous behavioral authentication pipeline showing client-side zero-content harvesting, dual-model hybrid inference, asymmetric EMA temporal filtering, and dynamic step-up TOTP security gating.*

A. Architectural Decomposition

The system comprises three decoupled operational tiers:

1) **Client Behavioral Harvester:** Runs inside the web browser via lightweight, passive DOM listeners. Computes interval timings and coordinate differences over sliding temporal windows without intercepting native input flow.
2) **Backend Hybrid ML Engine:** Implemented in Python FastAPI. Performs concurrent inference across the user-specific Isolation Forest and the statistical baseline profile.
3) **Adaptive Threat & Policy Engine:** Smooths risk using an asymmetric EMA filter, maintains audit logs with feature explainability, and enforces step-up TOTP challenges for high-value actions.

B. Zero-Content Privacy Proof

User privacy is guaranteed by design: When a key event occurs, the listener queries solely `event.timeStamp`. The properties `event.key`, `event.code`, and target input element values are never accessed or stored. Cursor listeners capture only coordinate displacements (Δx, Δy) and elapsed time (Δt); desktop absolute coordinates and external screen bounds are excluded. The telemetry payload consists strictly of aggregate statistical moments (μ, σ, ratios), rendering text reconstruction mathematically impossible.

IV. Mathematical Formulations & Feature Engineering

The behavioral extraction engine projects raw interactions into a compact 16-dimensional vector $\mathbf{x} \in \Re^{16}$.

TABLE I: 16-Dimensional Kinematic Feature Vector

#	Feature	Unit	Neuromuscular Description
1	mean_hold_time	ms	Average key depression duration
2	std_hold_time	ms	Variability in key depression
3	mean_flight_time	ms	Inter-key transition flight latency
4	std_flight_time	ms	Key transition flight variance
5	typing_speed_cpm	cpm	Interaction keystroke frequency
6	pause_ratio	ratio	Reflective hesitations (>400ms)
7	mean_velocity	px/ms	Average cursor movement speed
8	max_velocity	px/ms	Peak cursor ballistic speed burst
9	mean_acceleration	px/ms²	Rate of cursor velocity change
10	mean_jerk	px/ms³	Kinematic tremor / smoothness
11	curvature_ratio	ratio	Arc length / Euclidean chord
12	mean_click_duration	ms	Mouse button depression hold time
13	move_to_click_ratio	ratio	Navigational strokes per click
14	scroll_intensity	px/s	Wheel displacement speed/frequency
15	keystroke_overlap	ratio	Roll-typing finger overlap ratio
16	flight_time_cv	ratio	Cadence coefficient of variation

A. Keystroke Dynamics Modeling

Let a sequence of N consecutive keystrokes be denoted by timestamp tuples: $K = \{(t_{down,\,k}, t_{up,\,k})\}_{k=1}^{N}$.

1) Key Hold Duration (H_k):

$$H_k = t_{up,\,k} - t_{down,\,k} \tag{1}$$

The sample mean μ_H and standard deviation σ_H quantify key depression habits:

$$\mu_H = \frac{1}{N}\sum H_k, \quad \sigma_H = \frac{1}{N-1}\sum (H_k - \mu_H)^2 \tag{2}$$

2) Inter-Key Flight Latency (F_k) and Down-to-Down Digraph (D_k):

$$F_k = t_{down,\,k} - t_{up,\,k-1}, \quad D_k = t_{down,\,k} - t_{down,\,k-1} \tag{3}$$

3) Keystroke Overlap Ratio ($R_{overlap}$ / Roll-Typing Metric): Proficient touch-typists exhibit fluid "roll typing," where the next key is pressed before the previous key is released ($t_{down,\,k} < t_{up,\,k-1}$), yielding negative flight times. Hunt-and-peck typists release keys fully before pressing the next. The overlap ratio is formulated as:

$$R_{overlap} = \frac{1}{N-1}\sum_{k=2}^{N} \mathbb{I}(t_{down,\,k} < t_{up,\,k-1}) \tag{4}$$

where $\mathbb{I}(\cdot)$ is the indicator function.

4) Cadence Uniformity (CV_{flight}): Rhythmic consistency is captured by the dimensionless coefficient of variation:

$$CV_{flight} = \frac{\sigma_F}{\max(\mu_F, \varepsilon)}, \quad (\varepsilon = 1.0\ ms) \tag{5}$$

B. Mouse Kinematics Modeling

Let a trajectory be sampled at discrete time steps: $M = \{(x_i, y_i, t_i)\}_{i=1}^{M}$. Spatial displacement is $\Delta s_i = \sqrt{((x_i - x_{i-1})^2 + (y_i - y_{i-1})^2)}$, with velocity $v_i = \Delta s_i / \Delta t_i$ and acceleration $a_i = (v_i - v_{i-1}) / \Delta t_i$.

1) Kinematic Jerk (J_i / Motor Tremor): The third derivative of position with respect to time represents kinematic jerk:

$$J_i = \frac{d^3 s}{dt^3} = \frac{a_i - a_{i-1}}{\Delta t_i}, \quad \mu_J = \frac{1}{M-3}\sum_{i=4}^{M} |J_i| \tag{6}$$

Human motor control minimizes jerk to maximize trajectory smoothness [7]. Operating another individual's mouse or differing hand postures induces marked shifts in mean jerk.

2) Trajectory Curvature Ratio (C): Evaluates trajectory directness versus meandering:

$$C = \frac{\sum \Delta s_i}{\sqrt{((x_M - x_1)^2 + (y_M - y_1)^2)} + \varepsilon} \tag{7}$$

V. Hybrid Machine Learning Engine

A. RobustScaler Normalization

To prevent extreme outliers from distorting the training space, feature vectors are normalized using the median and Interquartile Range (IQR):

$$x'_j = \frac{x_j - median(X_j)}{Q_3(X_j) - Q_1(X_j)} \tag{8}$$

B. Model 1: Personalized Isolation Forest

Trained on K enrollment vectors ($K \geq 50$) with $T = 100$ trees, sub-sample $\psi = \min(256, K)$, and contamination $c = 0.05$. The anomaly score $s(\mathbf{x}, n)$ across path length $E(h(\mathbf{x}))$ is:

$$s(\mathbf{x}, n) = 2^{-\frac{E(h(\mathbf{x}))}{c(n)}} \tag{9}$$

where $c(n) = 2\ln(n-1) + 0.5772156649 - 2(n-1)/n$. The raw decision score $d(\mathbf{x})$ is mapped to percentage risk R_{IF} via calibrated sigmoid:

$$R_{IF} = \frac{100}{1 + e^{-18.0 \cdot (d(\mathbf{x}) - 0.04)}} \tag{10}$$

C. Model 2: Multivariate Parametric Z-Distance

To detect impostors whose metrics fall outside the enrolled convex hull without triggering tree boundaries, Model 2 computes standardized divergence against baseline profile $B = \{(\mu_j, \sigma_j)\}$. A safe variance floor prevents division by zero:

$$\sigma_{safe,\,j} = \max(\sigma_j, \ 0.08 \cdot \max(|\mu_j|, 1.0)) \tag{11}$$

Standardized deviations $Z_j = |x_j - \mu_j| / \sigma_{safe,\,j}$ are sorted descendingly. The mean of the top three most divergent features is evaluated:

$$Z_{top3} = \frac{1}{3}\sum_{k=1}^{3} Z_{sorted}[k] \tag{12}$$

Statistical risk R_{dist} is calibrated via logistic curve:

$$R_{dist} = \frac{100}{1 + e^{-2.4 \cdot (Z_{top3} - 1.90)}} \tag{13}$$

D. Hybrid Arbitrator Decision

The instantaneous risk score R_{inst} arbitrates between tree partitioning and parametric divergence:

$$R_{inst} = \min(100.0, \ \max(R_{IF}, \ 0.95 \cdot R_{dist})) \tag{14}$$

VI. Continuous Threat Scoring & Temporal Dynamics

A. Flaw of Memoryless Classifiers

Independent window classifiers exhibit two failure modes: 1) *False positive spikes:* A user momentarily fumbling a key trips immediate lockout. 2) *Pause evasion:* An impostor types a malicious command, pauses for 10 seconds to allow memoryless risk to decay to zero, and resumes attack in bursts.

B. Asymmetric Exponential Moving Average (EMA)

The session risk score S_t updates through an asymmetric EMA filter:

$$S_t = \alpha_{eff} \cdot R_{inst,\,t} + (1 - \alpha_{eff}) \cdot S_{t-1} \tag{15}$$

where α_{eff} is governed by a stateful Heaviside condition:

$$\alpha_{eff} = \begin{cases} 0.72 & \text{if } R_{inst,\,t} > S_{t-1} \quad \textit{(Threat Escalation)} \\ 0.18 & \text{if } R_{inst,\,t} \leq S_{t-1} \quad \textit{(Hysteresis Recovery)} \end{cases} \tag{16}$$

Security Rationale: Fast escalation ($\alpha = 0.72$) reaches critical lock threshold (>70%) in 1–2 consecutive anomaly windows (<2.5s). Slow decay ($\alpha = 0.18$) prevents adversaries from clearing elevated threat state by pausing interaction.

VII. Experimental Setup & Empirical Results

Evaluations were performed across 1,000 distinct observation windows comprising genuine enrolled sessions and physical impostor takeovers.

TABLE II: Confusion Matrix Evaluation (N=1,000 Windows)

Ground Truth	Predicted Genuine (Risk ≤ 40%)	Predicted Anomaly (Risk > 40%)
Genuine User (N=500)	480 (True Genuine - 96.0%)	20 (False Rejection - 4.0%)
Impostor User (N=500)	19 (False Acceptance - 3.8%)	481 (True Detected - 96.2%)

TABLE III: Comparative Model Performance

Model / Architecture	ROC-AUC	EER	FAR	FRR	F1-Score
One-Class SVM (RBF)	0.891	11.2%	10.8%	11.6%	0.884
Mahalanobis Z-Score Alone	0.934	7.1%	6.9%	7.3%	0.927
Isolation Forest Alone	0.952	5.4%	5.2%	5.6%	0.945
Proposed Hybrid Ensemble (EMA)	**0.982**	**3.9%**	**3.8%**	**4.1%**	**0.962**

A. Latency & Resource Utilization

• **Client Extractor:** <0.4% CPU overhead; window extraction takes 1.2 ms in browser.
• **Network Payload:** 412 bytes per evaluation window via gzip compression.
• **Backend Inference:** 8.4 ms mean latency (Isolation Forest: 6.1 ms, Z-distance: 2.3 ms).

VIII. Security Policy Enforcement & Gated Actions

The platform couples continuous risk scores directly with application authorization gates:

1) **Dynamic Sensitivity Slider ($\theta_{gate} \in [40\%, 90\%]$):** Configurable per organizational security posture (e.g., 50% for high-value banking wire transfers, 70% for standard dashboards).
2) **Step-Up MFA (RFC 6238 TOTP):** When session risk exceeds θ_{gate}, privileged actions return `HTTP 403 Forbidden` with `X-Risk-Challenge: MFA_REQUIRED`. The user must provide a 6-digit TOTP code to unlock the session.
3) **Emergency Recovery:** Single-use cryptographic backup codes prevent permanent lockouts in emergency scenarios.

IX. Threat Modeling & Adversarial Robustness

• **Physical Workstation Hijacking:** Intercepted within 2.1 seconds due to immediate shifts in key hold durations, roll-typing overlap, and cursor curvature.
• **Robotic Mouse Injections:** Synthetic script trajectories lack natural kinematic tremor (jerk $J \to 0$) and curve variation ($C = 1.000$). The statistical Z-distance flags unnatural uniformity ($Z_{jerk} > 4.5$), escalating risk to 100%.
• **Biological Fatigue Drift:** While typing speed may drop, intrinsic relative ratios (hold-to-flight, overlap ratio, path curvature) remain stable, buffered by *RobustScaler* and EMA hysteresis decay.

X. Limitations & Future Directions

Current limitations include a 90–120 second cold-start calibration period and sensitivity to extreme keyboard hardware transitions. Future research will investigate transfer learning across user population clusters to initialize zero-shot baseline profiles, and extend kinematic formulations to mobile touchscreen swipe dynamics and onboard IMU accelerometer sensors.

XI. Conclusion

Static point-in-time authentication protocols fail to defend against post-login session takeover, token exfiltration, and unattended terminal hijacking. This paper has presented, implemented, and empirically validated an AI-based continuous authentication system founded on neuromuscular behavioral biometrics. By combining a 16-dimensional zero-content kinematic feature vector, a dual-engine hybrid ensemble (Isolation Forest + parametric Z-distance), and an asymmetric temporal EMA threat filter, the system achieves sub-2.5 second threat response with an **ROC-AUC of 0.982**, **FAR of 3.8%**, and **FRR of 4.1%**, proving that continuous Zero-Trust session security can be achieved with zero user friction.

References

[1] F. Monrose and A. D. Rubin, "Keystroke dynamics as a biometric for authentication," *Future Generation Computer Systems*, vol. 16, no. 4, pp. 351–359, 2000.

[2] K. S. Killourhy and R. A. Maxion, "Comparing anomaly-detection algorithms for keystroke dynamics," in *Proc. IEEE/IFIP Int. Conf. on Dependable Systems & Networks (DSN)*, 2009, pp. 125–134.

[3] P. Gamboa and A. L. Fred, "A behavioral biometric system based on human computer interaction," in *Proc. SPIE Biometric Technology for Human Identification*, vol. 5404, 2004, pp. 381–392.

[4] C. Shen, Z. Cai, and X. Guan, "Continuous authentication for mouse dynamics using hidden Markov models," *IEEE Trans. Inf. Forensics Security*, vol. 8, no. 1, pp. 200–211, 2013.

[5] F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation Forest," in *Proc. 8th IEEE Int. Conf. Data Mining (ICDM)*, 2008, pp. 413–422.

[6] P. M. Fitts, "The information capacity of the human motor system in controlling the amplitude of movement," *J. Exp. Psychol.*, vol. 47, no. 6, pp. 381–391, 1954.

[7] T. Flash and N. Hogan, "The coordination of arm movements: an experimentally confirmed mathematical model," *J. Neurosci.*, vol. 5, no. 7, pp. 1688–1703, 1985.

[8] R. Joyce and G. Gupta, "Identity authentication based on keystroke latencies," *Commun. ACM*, vol. 33, no. 2, pp. 168–176, 1990.

[9] National Institute of Standards and Technology (NIST), "Zero Trust Architecture," *NIST SP 800-207*, Aug. 2020.

[10] D. M'Raihi, S. Machani, M. Pei, and J. Rydell, "TOTP: Time-Based One-Time Password Algorithm," *IETF RFC 6238*, May 2011.

[11] E. Frank, M. Hall, and I. Witten, *The WEKA Workbench: Data Mining Practical Machine Learning Tools and Techniques*, Morgan Kaufmann, 2016.

[12] M. E. Schuckers, *Computational Methods in Biometric Authentication*, Springer, 2010.